

Optimisation of Holistically-Nested Edge Detection: A Systematic Literature Review

Angel Mabuza^a, Mzoxolo Mbini^a, Thambo Nyathi^a

^a*Department of Computer Science, University of Pretoria, Pretoria, South Africa*

Abstract

Holistically Nested Edge Detection (HED) has become one of the most widely adopted deep learning architectures for boundary detection, formulating edge prediction as an image-to-image problem using a multiscale side-output mechanism. The ten years since its introduction have produced a large number of architectural variants, loss function modifications, and domain-specific adaptations. However, these contributions have emerged without systematic synthesis, leaving practitioners without evidence-based guidance for selecting architectures suited to specific accuracy, latency, and domain constraints requirements.

This systematic literature review (SLR) addresses this gap by synthesising the HED optimisation literature into a structured, reusable evidence base.

We incorporated the PICOC (Population, Intervention, Comparison, Outcome, Context) framework for scope definition and adhered to the PRISMA 2020 guidelines for reporting transparency. A total of 42 primary studies published between 2015 and 2025 were identified through six major academic databases and assessed using a 12-criterion weighted quality assessment instrument validated by a modified Delphi process.

No single optimisation strategy dominated all evaluation criteria. Backbone replacement (ResNet-50) increases computational cost but delivers consistent accuracy gains (median +1.8 percentage points Optimal Dataset Scale, ODS). Lightweight architectures (LDC, DexiNed) achieve 3–4× speedups with modest accuracy trade-offs. Attention mechanisms offer reliable but bounded gains (median +1.4 pp ODS) with minimal computational costs. Meta-heuristic hybrids

(Social Spider Optimisation, Black Widow Optimisation) show the highest reported improvement potential (median +2.6 pp ODS across three contributing studies on BSDS500); however, this claim requires cautious interpretation given the small number of contributing studies and the absence of statistical significance testing in source papers. Reproducibility is a critical concern: 61.9% of domain-specific studies rely on proprietary datasets that are not publicly available, and only 19% release their source code. The 12-criterion quality assessment yielded an average normalised score of 68.1% across all 42 studies, indicating acceptable but uneven methodological standards, with significant gaps in code availability and computational cost reporting.

A five-theme taxonomy hybrid architectural framework, multi-scale feature fusion, attention mechanisms, domain-specific optimisation, and meta-heuristic hybridisation characterise the field’s progression from generic accuracy improvements (2015–2018) toward deployment-aware, domain-aware paradigms (2022–2025). The taxonomy, quality assessment framework, and dataset mapping provided here are intended as reusable resources for advancing the reproducible and responsible deployment of edge-detection technology.

Keywords: Edge Detection, Holistically-Nested Networks, Systematic Literature Review, Deep Learning Optimisation, Meta-heuristics, Computer Vision

1. Introduction

Image segmentation is an invaluable technique in computer vision. Edge detection is a foundational step in this process, identifying regions in an image defined by specific criteria. The simplest approach is to locate points in a digital image where the brightness or intensity changes abruptly Xie & Tu (2015). These sharp transitions, or "discontinuities", usually correspond to object boundaries, surface textures, or changes in lighting and shadow.

Edge detection methods are broadly categorised into classical and deep learning approaches. The evolution from classical gradient-based operators to data-

driven approaches has substantially improved the robustness of boundary detection, particularly under challenging conditions, such as noise, texture variation, and illumination change. Convolutional Neural Networks (CNNs) now dominate this task, with Holistically Nested Edge Detection (HED) Xie & Tu (2015) emerging as a seminal architecture that formulates edge detection as an image-to-image prediction problem. HED’s key contribution is its multiscale side-output mechanism, which simultaneously captures fine-grained local details through shallow layers and semantic boundary information through deeper layers, producing strong results on benchmark datasets, including BSDS500 Arbeláez et al. (2011) and NYUD Silberman et al. (2012).

Despite their architectural elegance, standard HED models exhibit fundamental limitations that constrain their practical applications. The original implementation used a VGG-16 backbone with 138 million parameters, requiring 15.5 GFLOPs per forward pass. These computational demands preclude real-time inference on resource-constrained platforms, such as mobile devices, embedded systems, and edge computing nodes. HED is also sensitive to hyperparameter configuration, with performance degrading when learning rates, loss balancing coefficients or data augmentation strategies deviate from empirically determined defaults. These limitations are amplified in domain-specific contexts, such as heterogeneous intensity distributions in medical imaging (MRI, CT), extreme aspect ratios in remote sensing, and strict latency requirements in autonomous systems, all of which demand architectural modifications that the original formulation does not address.

Therefore, HED optimisation has emerged as a distinct and active research trajectory. These approaches include backbone replacement (ResNet, DenseNet, ConvNeXt), attention mechanism integration, loss function engineering, meta-heuristic hyperparameter optimisation, neural architecture search, and hybrid multi-task frameworks Heidler et al. (2021); Chen et al. (2024); Zhang et al. (2024). These techniques have proliferated without systematic synthesis, resulting in a fragmented knowledge base. Critical questions remain unresolved: Which optimisation strategies offer the best accuracy-efficiency trade-offs? Do

architectural modifications generalise across application domains? What reproducibility standards characterise the current research practice?

Existing surveys on edge detection Zhou et al. (2021); Acuna & Fidler (2022) provide broad taxonomies of deep learning approaches but lack a focused analysis of HED optimisation. Technical studies proposing HED variants Li et al. (2022); Melikyan & Melikian (2024) typically benchmark against baseline HED without contextualising their contributions within the broader optimisation landscape. Three specific gaps are identified.

First, there is **no comprehensive taxonomy** that organises HED optimisation techniques by architectural targets, computational mechanisms, and deployment constraints. Practitioners lack systematic guidance for selecting techniques based on specific accuracy, latency, and resource requirements.

Second, **quality assessment is absent**. Unlike related fields, where reproducibility standards have become normative Pineau et al. (2021), HED optimisation studies exhibit variable methodological rigor. Many do not report computational costs, conduct statistical significance testing, or discuss failure modes Hu & Saeli (2021); Li & Zhang (2016). Without critical appraisal, the field risks accumulating incremental benchmark improvements that fail to address the fundamental deployment constraints.

Third, **accuracy-efficiency trade-offs are poorly quantified**. Cross-study comparisons are hindered by inconsistent evaluation protocols, varying hardware specifications, and incomplete cost-reports. No meta-analytic synthesis has been conducted to identify Pareto-optimal configurations.

This review addresses these gaps through a systematic methodology, establishing a reproducible evidence base for meta-analytic comparisons and critical quality appraisals of HED optimisation research.

2. Related Surveys

Several surveys cover edge detection and boundary learning more broadly, and it is important to clarify how this study differs from and complements them.

Hu (2025) Hu (2025) presents a mathematical survey of deep edge detection from convolution to attention across multiple architectures, including HED, transformers (ECT, EdgeNAT), and generative models (DiffusionEdge). This survey covers a wide landscape but treats HED as one data point among many, rather than the focus of investigation. It does not provide a quality assessment of individual studies, quantify reproducibility failures across the HED literature, or synthesise domain-specific deployment constraints.

Sun et al. (2022) Sun (2022) survey traditional and deep learning edge detectors with backbone analysis and BSDS500/NYUD experiments. Their scope is similarly broad, covering HED alongside RCF, CASENet, and BDCN, without isolating the optimisation strategies specific to HED or assessing reproducibility standards across studies.

Mubashar et al. (2022) Mubashar (2022) analyse dataset limitations and deep neural network techniques in edge detection, referencing HED as a strong baseline. Their focus is on dataset characterisation rather than architectural optimisation of a specific model.

The present study makes four contributions that none of these surveys provide. First, this is the first SLR dedicated exclusively to HED optimisation, enabling a depth of analysis that is not possible in broader surveys. Second, it empirically quantifies reproducibility failures (61.9% proprietary datasets, 19% code release) using a structured quality assessment instrument. Third, it synthesises deployment-aware accuracy-latency-domain trade-offs across 42 studies. Fourth, it produces a reusable 12-criterion quality assessment instrument and taxonomy applicable to future HED and related CNN optimisation studies.

3. Research Methodology

This study follows a systematic literature review methodology to comprehensively analyse optimisation techniques for the Holistically-Nested Edge Detection (HED) algorithm. The review protocol integrates the PICOC framework for scope definition and the PRISMA 2020 guidelines for reporting trans-

parency. The methodological process is illustrated in Figure 1. All SLR materials search queries, screening forms, quality assessment scoring sheets, and the complete 42-study metadata list are publicly available at [https://github.com/\[repository-link-to-be-added-at-submission\]](https://github.com/[repository-link-to-be-added-at-submission]).




Figure 1: The methodology processes used for the study

3.1. Theoretical Framework: PICOC and PRISMA

The review protocol integrated two frameworks.

- **PICOC** Kitchenham & Charters (2007): Used to define eligibility criteria, construct the search string, and derive research questions specific to HED optimisation.
- **PRISMA 2020** Page et al. (2021): Used for reporting the systematic search, selection flow, and risk of bias assessment.

3.1.1. PICOC Framework Derivation

Table 1 operationalises the review scope across the five PICOC dimensions.

Table 1: PICOC Framework for HED Optimisation Scope Definition		
Element	Definition	Operationalisation in this Review
P	Population/Problem	Natural images (BSDS500, NYUD, PASCAL VOC); edge detection; thick vs. thin boundaries
I	Intervention	Meta-heuristics (PSO, GA, ABC, GWO); architectural modifications; post-processing; loss engineering; pruning/quantisation
C	Comparison	Original/baseline HED Xie & Tu (2015); optimisation algorithm comparisons; HED vs. optimised variants
O	Outcomes	ODS (Optimal Dataset Scale), OIS (Optimal Image Scale), AP (Average Precision); inference time (FPS); model size; boundary precision/recall
C	Context	General vision; medical imaging; industrial inspection; remote sensing; mobile/edge deployment

3.2. Research Questions

The research questions in Table 2 were formulated using the PICOC framework and provide a basis for coherent evidence synthesis.

Table 2: Research Questions and Motivation

RQ	Question	Motivation
RQ1	What optimisation techniques have been applied to HED architecture and training, and how do they compare in terms of accuracy-computation trade-offs?	To systematically catalogue architectural, algorithmic, and training interventions; to identify Pareto-optimal configurations for accuracy vs. inference speed vs. model size
RQ2	In which application domains has optimised HED been deployed, and what domain-specific constraints drive architectural choices?	To understand practical deployment contexts; to identify whether domain adaptations transfer across fields or require bespoke solutions
RQ3	Which preprocessing and postprocessing pipeline components are essential for HED optimisation performance?	To investigate algorithmic dependencies beyond network architecture; to identify standardised vs. domain-specific pipeline elements
RQ4	Which performance metrics are used, and how do optimised variants statistically compare with the baseline HED across benchmark datasets?	To assess quantitative improvement evidence; to evaluate metric diversity and reporting consistency; to enable meta-analytic comparison
RQ5	What are the major technical challenges, reproducibility gaps, and open research questions in HED optimisation?	To identify persistent limitations in current approaches; to prioritise future research directions based on evidence quality
RQ6	How do hybrid approaches (multi-task learning, multi-objective optimisation, and meta-heuristic-neuro hybridisation) advance beyond single-objective architectural optimisation?	To explore advanced optimisation paradigms that transcend individual technique categories; to assess synergistic vs. additive benefits

3.3. Search Strategy and Keywords

An electronic search was conducted in six major academic databases, IEE-EXplore, Engineering Village, ScienceDirect, Scopus, SpringerLink, and ACM Digital Library, covering publications from 2015 (the year of HED’s introduction) to 2025. Search strings were derived from the PICOC framework as follows:

```
("holistic nested edge detection" OR "HED" OR "holistic nested network")
AND ("optimisation" OR "optimization" OR "tuning" OR "hyperparameter"
OR "enhance" OR "improve")
```



```

AND ("genetic algorithm*" OR "particle swarm" OR "PSO"
    OR "differential evolution"
    OR "ant colony" OR "artificial bee colony" OR "grey wolf"
    OR "metaheuristic*" OR "deep learning" OR "neural network"
    OR "CNN" or "VGGNet")
OR ("architectural modification" OR "attention mechanism"
    OR "dilated convolution")
OR ("post-processing" OR "conditional random field" OR "CRF" OR "NMS")
AND ("edge detection" OR "boundary detection" OR "contours")
AND ("ODS" OR "OIS" OR "average precision")
AND ("image segmentation" OR "segmentation" OR "image classification"
    OR "classification")

```

3.3.1. Search strategy per database

Table 3 lists the exact search strings used for each database, as the search functionality varies across platforms.

Table 3: Search Strings Used Across Databases for HED Network Optimisation Literature

Database	Search String
IEEEExplore	<pre> ("holistic nested edge detection" OR "HED network" OR "holistic nested" OR "HED") AND ("optimiz*" OR "tun*" OR "enhanc*" OR "improv*") </pre>
Engineering Village	<pre> (((((((holistic nested edge detection OR HED network OR holistically-nested OR HED)) WN ALL) AND ((optimiz* OR improv* OR enhanc* OR tun*) WN ALL))) AND ((2025 OR 2024 OR 2023 OR 2022 OR 2021 OR 2020 OR 2019 OR 2018 OR 2017 OR 2016 OR 2015) WN YR)) AND (({ja} OR {ca}) WN DT)) AND ({english} WN LA) </pre>

Database	Search String
ScienceDirect, Scopus, SpringerLink, ACM Digital Library, Web of Science	("holistic nested edge detection" OR "HED") AND (optimize OR optimized OR optimization OR enhance OR enhanced OR improvement)

3.4. Inclusion and Exclusion Criteria

Studies were included if they (1) proposed, evaluated, or systematically compared an optimisation technique applied to the HED architecture or a direct HED derivative; (2) reported at least one quantitative performance metric (ODS, OIS, AP, AUC, F1, or inference time); (3) were published in English in a peer-reviewed journal or conference proceedings between 2015 and 2025; and (4) were retrievable in full text.

Studies were excluded if they: (1) used HED purely as a component without evaluating or modifying its optimisation; (2) were literature reviews, surveys, or secondary studies; (3) appeared only as preprints without peer review at the time of screening; (4) reported results solely on proprietary datasets with no comparison to a standardised baseline; or (5) were duplicate publications of the same experimental work.

3.5. Screening Process and Inter-Reviewer Agreement

The screening was conducted in two stages. In Stage 1, two reviewers independently screened all 42,368 retrieved titles and abstracts against the inclusion/exclusion criteria, resolving disagreements through discussion. In Stage 2, the full texts of all candidate studies were independently assessed by the same two reviewers. A third reviewer resolved any remaining disagreements. Inter-reviewer agreement at Stage 2 was calculated using Cohen’s κ statistic, yielding $\kappa = 0.81$, which indicates a strong agreement Landis & Koch (1977). The complete screening flow is shown in Figure 2.

3.6. Search Results and Screening

A total of 42,368 articles were identified in the six databases. After deduplication and two-stage screening, 42 primary studies were included in the review. The complete PRISMA 2020 flow diagram is presented in Figure 2.




Figure 2: PRISMA 2020 search and screening flow showing per-database retrieval counts, deduplication, title/abstract screening, full-text eligibility assessment, and final inclusion. The inclusion and exclusion counts at each stage followed the PRISMA 2020 reporting standard Page et al. (2021).

3.7. Quality Assessment

A 12-criterion Quality Assessment Criteria (QAC) instrument was adapted for HED optimisation studies (Table 4). The average normalised quality score across the 42 included studies was 68.1%, which we interpret as acceptable but uneven, comparable to the 65–70% benchmarks reported in computer science SLRs of similar scope Kitchenham & Charters (2007). Four studies (9.5%) Shoukat et al. (2023), Li & Zhang (2016), Yu et al. (2022), and Zhong et al. (2023) fell below the 50% threshold. These are retained as weak-evidence entries to preserve domain coverage but are excluded from the quantitative comparisons in Section 5.

Inter-rater reliability for QAC scoring was assessed across all 42 studies using the intraclass correlation coefficient (ICC), yielding $ICC = 0.79$ (95% CI: 0.71–0.86), indicating good reliability Koo & Li (2016). Disagreements exceeding one Likert point were resolved through structured discussions among all three authors.

Each criterion was assigned a weight (High, Medium, or Low) through a modified Delphi process Dalkey & Helmer (1963) involving three rounds of expert consultation among the authors. Table 5 documents this rationale.

Table 5: Rationale for Quality Assessment Weight Assignments

Weight	Criteria	Justification
High (3×)	Q2, Q3, Q4, Q7, Q11	Core reproducibility and validity requirements: the methodology must be well-defined (Q2), compared against a baseline (Q3), with implementation details documented (Q4), appropriate metrics used (Q7), and code available for verification (Q11). Without these, the findings cannot be independently validated or built upon.
Medium (2×)	Q1, Q5, Q6, Q8, Q9, Q12	Important for contextual interpretation: problem definition (Q1), pipeline documentation (Q5), cross-dataset robustness (Q6), computational cost reporting (Q8), conclusion validity (Q9), and domain constraint addressing (Q12) enable proper assessment of generalisability and practical utility.
Low (1×)	Q10	Valuable but supplementary: discussion of limitations indicates research maturity but does not directly affect the evidence quality for the reported optimisation technique.

Table 4: Quality Assessment Criteria (QAC) with definitions and scoring guidance

Q	Quality Assessment Question	Definition / Scoring Guidance	Weight
Q1	Is the optimisation problem or HED modification clearly defined?	Score 5 if the research objective, target component, and expected improvement are unambiguously stated. A score of 1 was assigned if the problem was vague or unstated.	Medium
Q2	Is the proposed optimisation methodology well-defined and reproducible?	Score 5 if all algorithmic steps, hyperparameter settings, and training procedures are documented in sufficient detail to replicate. A score of 1 was assigned if the methodology was described at a high level only.	High
Q3	Did the study compare baseline HED using standard metrics?	Score 5 if results are reported against the original HED Xie & Tu (2015) on a standard benchmark (BSDS500, NYUD, PROMISE12, etc.). A score of 1 was assigned if no baseline comparison existed.	High
Q4	Are the implementation details documented?	Score 5 if the software framework, hardware, training schedule, and data splits are all reported. A score of 1 was assigned if the implementation details were absent.	High
Q5	Are the preprocessing/postprocessing steps documented?	Score 5 if the full data pipeline (augmentation, normalisation, filtering, thinning) is described. A score of 1 was assigned if the pipeline was not mentioned.	Medium
Q6	Did the study evaluate robustness across multiple datasets?	Score 5 if the results are reported on two or more independent datasets. A score of 1 was assigned if only one dataset was used without cross-dataset validation.	Medium
Q7	Are the edge detection metrics reported appropriately?	Score 5 if ODS, OIS, AP, or domain-equivalent metrics were reported with clear definitions. A score of 1 was assigned if the metrics were undefined or absent.	High
Q8	Is the computational cost reported along with the accuracy?	Score 5 if FLOPs, parameters, inference time, and/or memory were reported on the specified hardware. A score of 1 was assigned if the cost was not mentioned.	Medium
Q9	Do the conclusions align with the empirical results?	Score 5 if the claims in the conclusion are directly supported by tabulated or graphed results. A score of 1 was assigned if the conclusions overstated or contradicted the results.	Medium
Q10	Are limitations and failure cases discussed?	Score 5 if failure modes, boundary conditions, and scope limitations are explicitly discussed. A score of 1 was assigned if limitations were absent.	Low
Q11	Is the source code or configuration publicly available?	Score 5 if a functional public repository (GitHub) was provided. A score of 1 was assigned if no code was available.	High
Q12	Does the study address domain-specific constraints?	Score 5 if constraints specific to the application domain (intensity distribution, class imbalance, and latency) are identified and addressed. A score of 1 was assigned if the domain context was ignored.	Medium

The numerical multipliers (high = 3, medium = 2, and low = 1) reflect ordinal importance scaling. A 1–5 Likert scale (1 = poor, 2 = below average, 3 = satisfactory, 4 = good, 5 = excellent) was used for each criterion. The weighted score was calculated as follows:

$$\text{Weighted Score}_i = \sum_{j=1}^{12} S_{ij} \times W_j$$

where S_{ij} is the score (1–5) for paper i on criterion j , and $W_j \in \{1, 2, 3\}$ is the weight multiplier. The normalised percentage is $\frac{\text{Weighted Score}_i}{140} \times 100$, where 140 is the maximum possible score (5 points $\times$ 46 total weight units). The 50% threshold is consistent with comparable computer science SLRs Kitchenham & Charters (2007).

Table 6 presents the complete quality assessment output. The distribution shows that 14 papers (33.3%) scored 70–79%, 13 papers (31.0%) scored 60–69%, 10 papers (23.8%) scored 80–89% (Good), and only one paper (2.4%) achieved 90% or above (Excellent). Four papers (9.5%) fell below the 50% threshold. The 68.1% mean reflects acceptable methodological standards, with persistent gaps in code availability and computational cost reporting.

Table 6: Quality Assessment Results

Paper	Q1	Q2	Q3	Q4	Q5	Q6	Q7	Q8	Q9	Q10	Q11	Q12	TWS	Norm.
Guo et al. (2023)	5	5	5	5	5	4	5	3	5	4	5	5	128	91.4
Cheng et al. (2017)	5	5	5	5	5	5	5	3	5	5	1	5	124	88.6
García-Peraza-Herrera et al. (2017)	5	5	5	5	3	4	5	5	5	3	1	5	120	85.7
Jain & Dwivedi (2024)	5	5	5	5	5	5	5	2	5	2	1	5	120	85.7
Sumathipala et al. (2018)	5	5	4	5	5	3	5	2	5	3	1	5	113	80.7
Hou et al. (2016) SOD	5	5	4	4	4	5	5	3	5	2	3	4	116	82.9
Soria et al. (2020)	5	5	5	4	5	5	5	3	5	3	3	4	122	87.1
Heidler et al. (2021)	5	5	4	4	4	4	5	3	5	2	3	4	110	78.6
Tseng & Sun (2023)	4	4	4	4	4	3	4	2	4	2	2	4	98	70.0
Chen et al. (2024)	4	4	4	4	4	3	4	2	4	2	1	4	94	67.1
Mulay et al. (2023)	4	4	4	4	4	3	4	2	4	2	1	4	92	65.7
Sreelakshmi & Vardhan (2024)	4	4	4	3	4	3	4	2	4	2	1	4	92	65.7
Li & He (2022)	4	4	3	3	3	3	4	2	4	2	1	3	82	58.6
Lou & Zang (2023)	4	4	3	3	3	3	4	2	4	2	1	3	82	58.6
Wang et al. (2018)	4	4	3	3	3	3	4	2	4	2	1	3	82	58.6

Continues on the next page...

Table 6 – continued from the previous page

Paper	Q1	Q2	Q3	Q4	Q5	Q6	Q7	Q8	Q9	Q10	Q11	Q12	TWS	Norm.
Ji et al. (2021)	4	4	4	4	4	4	4	2	4	2	1	4	98	70.0
Soria & Sappa (2021)	4	4	4	4	4	4	4	2	4	2	2	4	100	71.4
Jiao et al. (2022)	4	4	4	4	4	4	4	2	4	2	1	4	98	70.0
Xu et al. (2023)	4	4	4	4	4	4	4	2	4	2	1	4	98	70.0
Ji et al. (2023)	4	4	3	3	3	3	4	2	4	2	1	3	82	58.6
Nogues et al. (2016)	5	5	4	4	4	4	5	3	5	3	1	5	113	80.7
Beaini et al. (2020)	4	4	4	3	4	4	4	2	4	2	2	4	98	70.0
Ji et al. (2020)	5	5	4	5	5	3	5	2	5	3	1	5	113	80.7
Wang et al. (2022)	4	4	3	3	3	3	4	2	4	2	1	3	82	58.6
Hu & Saeli (2021)	4	4	3	3	3	3	3	2	3	2	1	3	78	55.7
Melikyan & Melikian (2024)	4	4	4	4	4	3	4	3	4	3	2	4	100	71.4
Yang et al. (2022)	4	4	4	3	4	3	4	2	4	2	1	4	92	65.7
Yuan et al. (2023)	4	4	3	3	3	3	4	2	4	2	1	3	82	58.6
Tiwari et al. (2024)	4	4	4	3	4	3	4	2	4	2	1	4	92	65.7
Zhao et al. (2023)	4	4	3	3	3	3	4	3	4	2	2	4	88	62.9
Gamal et al. (2023)	4	4	4	3	4	4	4	2	4	2	1	4	96	68.6
Jung et al. (2021)	4	4	4	4	4	4	4	2	4	2	1	4	98	70.0
Chen et al. (2018)	4	4	3	3	3	3	4	2	4	2	1	3	82	58.6
Li et al. (2021)	4	4	3	3	3	3	4	3	4	2	1	4	86	61.4
Zhang et al. (2024)	4	4	4	3	4	3	4	2	4	2	1	3	90	64.3
Li et al. (2022)	4	4	4	4	4	3	4	2	4	2	1	4	96	68.6
Gui et al. (2023)	4	4	4	4	4	3	4	2	4	2	1	4	96	68.6
Xu et al. (2022)	4	4	4	4	4	3	4	2	4	2	2	4	98	70.0
Wen et al. (2017)	4	4	4	4	4	3	4	3	4	2	1	4	98	70.0
Shoukat et al. (2023)	3	3	3	2	3	2	3	1	3	1	1	3	64	45.7
Li & Zhang (2016)	3	3	3	2	3	2	3	2	3	1	1	3	66	47.1
Yu et al. (2022)	3	3	3	3	3	2	3	2	3	1	1	3	68	48.6
Zhong et al. (2023)	3	3	3	3	3	2	3	2	3	1	1	3	68	48.6

4. Standard Holistically-Nested Edge Detection Algorithm

Holistically Nested Edge Detection (HED) represents a significant advancement in edge detection, formulating it as an image-to-image prediction task using deep convolutional neural networks. The standard architecture introduced by Xie and Tu Xie & Tu (2015) uses the VGG-16 network as its backbone, which is modified to produce edge maps at multiple scales.

4.1. Architectural Overview

The HED network generates side outputs at five layers (conv1_2, conv2_2, conv3_3, conv4_3, conv5_3), corresponding to feature maps with increasing receptive fields and decreasing spatial resolutions. Each side output produces an edge prediction map that is upsampled to the original image size. A weighted fusion layer combines these multiscale predictions to generate the final edge map.

Mathematically, given an input image I , HED produces side outputs $S^{(k)}$ for $k = 1, \dots, 5$ and a fused output F . The loss function incorporates both the side output losses and fusion loss:

$$\mathcal{L}_{\text{total}} = \sum_{k=1}^5 \alpha_k \mathcal{L}_{\text{side}}^{(k)} + \mathcal{L}_{\text{fuse}}$$

where α_k is the weight trained for each output.

4.2. Limitations of Standard HED

Despite its advantages, standard HED exhibit several limitations that motivate the optimisation strategies reviewed in this study.

1. **Computational complexity:** The VGG-16 backbone contains approximately 138 million parameters, resulting in high memory usage and inference times unsuitable for real-time applications.
2. **Hyperparameter sensitivity:** Performance is sensitive to learning rate, weight initialisation, and loss balancing parameters, requiring extensive manual tuning.
3. **Limited edge types:** While effective for prominent edges, the architecture struggles with fine textures, low-contrast boundaries, and complex junctions.
4. **Fixed architecture:** The predetermined side output locations and fusion strategy may not be optimal across all detection scenarios.
5. **Domain specificity:** Performance degrades when applied to domains significantly different from the training distribution.

4.3. Computational Complexity Analysis

For an input image of size $H \times W$, the forward pass through VGG-16 requires approximately 15.5×10^9 floating-point operations. The peak memory usage occurs at the conv3_3 layer, requiring storage for $256 \times (H/4) \times (W/4)$ activations. These

characteristics highlight the need for optimisation, particularly in resource-constrained or real-time settings.

4.3.1. Comparative Complexity Analysis of Optimised Variants

Table 7 summarises the computational metrics for the representative optimised variants. Only six of the 42 studies explicitly reported computational statistics. Backbone replacement with ResNet-50 increases the parameters by 12% but improves the gradient flow Yang et al. (2022). Lightweight Dense CNNs (LDC) reduce parameters by 89% with modest accuracy degradation Melikyan & Melikian (2024). Attention mechanisms add minimal parameters ($< 1\%$) but increase memory access patterns, affecting latency more than throughput.

Table 7: Computational Complexity Comparison: Standard HED vs. Optimised Variants

Variant	Backbone	Parameters	FLOPs	Memory (MB)	Relative Speed
Standard HED Xie & Tu (2015)	VGG-16	138M	15.5G	528	1.0×
ResNet-HED Yang et al. (2022)	ResNet-50	154M	18.2G	612	0.85×
LDC-HED Melikyan & Melikian (2024)	Lightweight Dense CNN	15M	4.2G	89	3.2×
DexiNed Soria et al. (2020)	DenseNet-inspired	9M	2.8G	64	4.5×
HED-UNet Heidler et al. (2021)	VGG-16 + UNet decoder	142M	16.1G	545	0.92×
Attention-HED Xu et al. (2023)	VGG-16 + CAM	138.5M	15.6G	534	0.95×

These metrics reveal a fundamental *accuracy-efficiency paradox*: modifications that improve boundary precision (ResNet, attention) typically increase computational cost, whereas efficiency-focused approaches (LDC, DexiNed) sacrifice the richness of the multi-scale features. This motivates the hybrid optimisation strategies discussed in Section 5 (Theme 5).

5. Taxonomy of HED Optimisation Variants

The following five themes were derived inductively from the 42 primary studies. Starting from individual study descriptions, we grouped the papers by the primary mechanism through which they modified or extended HED. The themes were iterated across two coding rounds by two independent authors, with disagreements resolved by the third author. Theme 1 (Hybrid Architectural Frameworks) covers models that structurally combine HED with a separate segmentation or detection backbone architecture. Theme 2 (Multi-Scale Feature Fusion) covers studies that innovate the aggregation of HED’s hierarchical side outputs without replacing the backbone. The boundary between these two themes is the presence of a secondary backbone: Theme

1 papers use two networks in parallel or in series, whereas Theme 2 papers modify the fusion mechanism within a single HED network. The distribution of studies across themes was as follows: Theme 1 (10 studies, 23.8%), Theme 2 (9 studies, 21.4%), Theme 3 (8 studies, 19.0%), Theme 4 (12 studies, 28.6%), and Theme 5 (3 studies, 7.1%). These numbers reflect each study’s primary theme; some studies contribute to more than one.

5.1. Theme 1: Hybrid Architectural Frameworks

Researchers in this theme combine edge detection with semantic segmentation or object detection in a single model. This addresses the limitations of each method used in isolation, producing more context-aware boundary detection results.

The SoUNet model integrates a side-output module with a batch normalisation layer into the U-Net structure for pavement crack identification on the CrackForest dataset Li et al. (2022). HED-UNet performs simultaneous segmentation and edge detection, outperforming single-task models for Antarctic coastline monitoring using in-house remote sensing imagery Tseng & Sun (2023).

Another hybrid approach combines an improved HED network with PSP-Net using a semantic and edge feature fusion module to handle small object disappearance in complex scenes. For liver segmentation, HED with Mask R-CNN produces high-precision primal sketches across different imaging modalities on the LiTS dataset Muly et al. (2023). Real-time 3D eyelid tracking uses a modified holistically nested network that jointly detects and identifies four distinct semantic eyelid edges from in-house video sequences Wen et al. (2017).

5.2. Theme 2: Multi-Scale and Hierarchical Feature Fusion

This theme focuses on how features from different CNN layers are combined to produce crisper boundaries and suppress noise. One approach introduces short connections between shallower and deeper side-output layers, giving activations the dual capability of highlighting objects while accurately locating boundaries, which is evaluated on the BSDS500, DUT-OMRON, and MSRA-B datasets Hou et al. (2016).

The Scale-Invariant Salient Edge Detection (SISED) framework uses a normalised Hadamard product as an AND gating operation to acknowledge saliency only when supported across multiple scales, benchmarked on BSDS500 and NYUD Hu & Saeli

(2021). Parallel CNN structures process individual colour channels (R, G, B) separately before fusing them into an edge approximation matrix, which is evaluated on BSDS500 alongside meta-heuristic refinement Zhang et al. (2024). The RCF (Richer Convolutional Features) network builds on HED by embedding finer details directly into the side-output feature maps through extra convolutional layers Hu & Saeli (2021).

5.3. Theme 3: Integration of Attention Mechanisms

Attention modules dynamically weight feature importance, which is crucial in complex backgrounds. Attentional Feature Fusion (AFF) allows the dynamic fusion of semantic and edge information, improving a model’s response to specific semantic scenes, as evaluated on BSDS500, PASCAL VOC, and ADE20K Wang et al. (2018).

Channel Attention Modules (CAM) were integrated at the end of each VGG-16 stage to selectively amplify channels with valid crack information while suppressing noise, evaluated on Bridge_Crack_Image_Data and CFD Xu et al. (2022). HA U-Net incorporates attention to ensure that lateral outputs at each hierarchical level retain both fine spatial detail and high-level semantic context, evaluated on ISPRS Potsdam, Vaihingen, and INRIA building datasets Xu et al. (2023).

5.4. Theme 4: Domain-Specific Optimisation

HED has been adapted to the distinct constraints of civil infrastructure, medicine, and remote sensing Li et al. (2022).

Pavement and Concrete Infrastructure: PCDM-HED uses Profile Component Decomposition to handle heterogeneous 3D data across seven in-house laser-scanning datasets Gui et al. (2023). CDFFHNet focuses on the linear topological structures of cracks in complex textures on Bridge_Crack_Image_Data and CFD Xu et al. (2022).

Medical Imaging: Research addresses heterogeneous intensity distributions in MRI and CT scans, including prostate segmentation on PROMISE12 Cheng et al. (2017), retinal vessel detection on DRIVE/STARE/CHASEDB1 Jain & Dwivedi (2024), and osteoarthritis grading on the OAI knee X-ray dataset Yang et al. (2022).

Remote Sensing and Civilian Infrastructure Monitoring: Sea-sky line detection fuses HED with ConvNeXt and Canny operators for varied marine environments using in-house maritime horizon data Chen et al. (2024). Urban infrastructure mapping and disaster response assessment use the DOTA and NWPU VHR-10 remote sensing

datasets to classify objects in aerial imagery for civilian applications, including building damage assessment and transportation network monitoring Tiwari et al. (2024).

Industrial and Agricultural: Studies extract plant root contours on edge devices using adaptive thresholding on in-house plant root imagery Zhao et al. (2023) and measure industrial workpiece dimensions online using optimised HED on in-house manufacturing datasets Li et al. (2021).

5.5. Theme 5: Fusion with Traditional ML and Meta-Heuristic Optimisation

Studies on this theme use HED as a feature extraction component within broader pipelines that include conventional machine-learning classifiers or optimisation methods. This is the smallest theme (three studies), and the results should be interpreted with caution, given the limited evidence base.

Feature Selection: The minimum redundancy maximum relevance (mRMR) algorithm selects the most relevant deep features from a hybrid PSP-Net/HED model before classification by XGBoost, evaluated on the DOTA and NWPU VHR-10 datasets for civilian infrastructure classification Tiwari et al. (2024).

Meta-heuristic optimisation: Black Widow Optimisation (BWO) and Social Spider Optimisation (SSO) refine feature sets and improve edges in colour images benchmarked on BSDS500 Zhang et al. (2024).

Hybrid Classification: For automated colon biopsy grading, HED-based segmentation is combined with LSTM and Random Forest classifiers through a majority voting framework on in-house histology slides Sreelakshmi & Vardhan (2024).

Operator Assistance: High-precision edge algorithms combine HED with an improved Sobel operator to enhance target localisation accuracy in aerial imagery for civilian applications Zhong et al. (2023).

6. System Designer’s Selection Guide

Practitioners face three interrelated choices: the backbone/architecture, preprocessing pipeline, and postprocessing configuration. Table 8 maps representative real-world deployment scenarios to optimal configurations based on the synthesised evidence. This is intended as a first-pass decision aid and not a deterministic prescription. For any given deployment, actual hardware profiling and dataset-specific ablation remain necessary.

Table 8: Expert System Selection Guide: Application Scenario to Pipeline Configuration

Scenario	Recommended Backbone	Preprocessing	Postprocessing	Evidence Source
High-throughput industrial surface inspection (latency <50 ms)	LDC or DexiNed (lightweight)	Otsu-Canny for rough edge seeding; normalise to [0,1]	NMS + double-thresholding; consider INT8 quantisation (<1% accuracy loss)	Melikyan & Melikian (2024); Soria et al. (2020)
High-precision medical scanning (MRI/CT; accuracy paramount)	ResNet-50 backbone with multi-scale fusion	IQR scaling; median filter to suppress acquisition noise	Skeletonisation; class-balanced Dice loss during training	Cheng et al. (2017); Yang et al. (2022)
Remote sensing / urban infrastructure mapping (resolution variability)	HED-UNet or HED + ConvNeXt	PCDM for 3D data; standard normalisation for 2D imagery	mRMR feature selection before classification; NMS	Heidler et al. (2021); Chen et al. (2024); Tiwari et al. (2024)
Edge device deployment (embedded, <15M parameters)	DexiNed (9M parameters) or LDC (15M)	Lightweight augmentation (horizontal flip, brightness jitter)	Fixed double-threshold; avoid meta-heuristic postprocessing (too slow on-device)	Soria et al. (2020); Melikyan & Melikian (2024)
Infrastructure crack detection (2D pavement, CFD/CrackForest)	CDFHNet or SoUNet with VGG-16 backbone	Otsu-Canny; histogram equalisation for low-contrast surfaces	NMS + morphological thinning	Li et al. (2022); Xu et al. (2022)
General computer vision benchmark (BSDS500, maximise ODS)	VGG-16 + attention (CAM or AFF)	Standard ImageNet normalisation	NMS; avoid aggressive postprocessing that trades recall for precision	Jiao et al. (2022); Xu et al. (2023)

For a concise architecture selection decision path, if the latency is less than 50 ms on the target hardware, DexiNed or LDC-HED should be preferred. If latency is not the primary constraint and the domain is medical or remote sensing, ResNet-50 with domain-specific preprocessing (IQR scaling for MRI, PCDM for 3D pavement) should

be used. If the goal is maximum ODS on the BSDS500, a channel or spatial attention module is added to the standard VGG-16 backbone. If postprocessing time is available (offline or batch processing), meta-heuristic refinement (SSO, BWO) offers additional ODS gains on colour images. **[AUTHOR ACTION REQUIRED: Convert this decision path into a formal flowchart figure for submission.]**

7. Systematic Mapping of the Results

7.1. RQ1: Optimisation Techniques and Accuracy-Computation Trade-offs

HED optimisation primarily targets the backbone architecture, loss functions, and feature fusion strategies.

For architectural interventions, researchers have replaced VGG-16 with ResNet-50 to mitigate vanishing gradients on the OAI knee X-ray dataset Yang et al. (2022) or with ConvNeXt for small-image marine environments Zhang et al. (2024). Other interventions include inserting squeeze-and-excitation (SE) modules to model channel relationships Li & He (2022), using dilated convolutions to expand receptive fields without resolution loss Li & He (2022), and introducing short connections between hierarchical layers to transfer semantic knowledge Hou et al. (2016).

In kernel and feature optimisation, larger convolutional kernels (5×5 or 7×7) improve the local image context for complex shapes, such as tumours in multiparametric prostate MRI, although they double arithmetic operations and risk overfitting Sumathipala et al. (2018).

Standard convolutions have been replaced with Pixel Difference Convolution (PDC) to improve the gradient sensitivity in noisy SEM images on BIPED and BSDS500 Melikyan & Melikian (2024). Residual variables (deformable) and dilated convolutions expand the receptive fields and adapt to diverse edge shapes without adding parameters Wang et al. (2022).

For loss function improvements, models now use class-balanced Dice loss and SSIM loss to handle pixel imbalance in medical scans Yang et al. (2022). Some frameworks are optimised solely on the fused output rather than the individual side outputs to reduce blurring in industrial workpiece measurements Li et al. (2021).

These interventions generally produce crisper edges than the baseline HED, which is cited for producing "rough" or "fuzzy" boundaries Wang et al. (2022). Backbone

replacements, such as ResNet-50, improve accuracy but increase the computational load, whereas lightweight modules, such as LDC, prioritise speed Melikyan & Melikian (2024).

7.2. RQ2: Application Domains and Domain-Specific Constraints

Optimised HED variants are deployed across sectors that require high-precision boundary delineation. Table 9 maps all 42 studies to their datasets, application domains, and key characteristics. BSDS500 Arbeláez et al. (2011) and NYUD Silberman et al. (2012) dominate general computer vision evaluation (appearing in 19 and 9 studies, respectively), while 61.9% of domain-specific studies rely on proprietary in-house datasets, which is the primary reproducibility concern identified in this review.

Table 9: Dataset Usage Across Included Studies by Application Domain

Study	Dataset(s) Used	Domain	Notes
Xie & Tu (2015)	BSDS500, NYUD, PASCAL VOC	General CV	Original HED baseline
Hou et al. (2016)	BSDS500, DUT-OMRON, MSRA-B, SOD	General CV / Saliency	Short connections; salient object detection
Soria et al. (2020)	BIPED, BSDS500, NYUD, MDBD, CID, DCD	General CV	DexiNed; six-dataset evaluation
Hu & Saeli (2021)	BSDS500, NYUD	General CV	Hadamard AND-gating across scales
Soria & Sappa (2021)	BSDS500, NYUD, RGB-NIR Scene Dataset	Multi-spectral	RGB + Near-infrared channel fusion
Wang et al. (2022)	BSDS500	General CV	Multi-feature side-output fusion
Li & He (2022)	BSDS500	General CV	SE modules and dilated convolutions
Lou & Zang (2023)	BSDS500	General CV	Incremental HED architectural study
Jiao et al. (2022)	BSDS500, NYUD, VOC2012	General CV	Attentional Feature Fusion (AFF)
Zhang et al. (2024)	BSDS500, in-house colour images	General CV	Parallel CNNs + SSO/BWO meta-heuristic
Yuan et al. (2023)	BSDS500, in-house aerial imagery	General CV	HED + improved Sobel operator
Ji et al. (2021)	BSDS500, PASCAL VOC	General CV	Parallel FCN per colour channel

Study	Dataset(s) Used	Domain	Notes
Cheng et al. (2017)	PROMISE12, in-house MRI	Medical	Prostate MRI segmentation with HNN
Sumathipala et al. (2018)	In-house multi-parametric MRI	Medical	Multi-institution prostate cancer detection
Ji et al. (2020)	PROMISE12, in-house MRI	Medical	Enhanced HNN for prostate segmentation
Nogues et al. (2016)	TCIA Lymph Node CT (NIH)	Medical	Lymph node cluster segmentation
Yang et al. (2022)	OAI (knee X-ray), in-house X-ray	Medical	MED-HED; osteoarthritis grading
Mulay et al. (2023)	LiTS, in-house CT/MRI	Medical	Liver segmentation; Mask R-CNN hybrid
Sreelakshmi & Vardhan (2024)	In-house colon histology slides	Medical	Biopsy grading; LSTM + Random Forest
Jain & Dwivedi (2024)	DRIVE, STARE, CHASEDB1	Medical	Retinal vessel detection
Gamal et al. (2023)	DMR-IR (thermography), in-house	Medical	Early breast cancer screening
Wen et al. (2017)	In-house eyelid video sequences	Medical / HCI	Real-time 3D eyelid tracking
Beaini et al. (2020)	In-house plant leaf imagery	Agricultural	Plant leaf disease saliency detection
Heidler et al. (2021)	In-house Antarctic coastline, BSDS500	Remote Sensing	HED-UNet; simultaneous seg. + edge
Tseng & Sun (2023)	In-house sea-land boundary dataset	Remote Sensing	Sea-land segmentation monitoring
Chen et al. (2024)	In-house maritime horizon dataset	Maritime	Sea-sky line; HED + ConvNeXt + Canny
Jung et al. (2021)	ISPRS Potsdam, SpaceNet, WHU Building	Remote Sensing	Building boundary extraction
Tiwari et al. (2024)	DOTA, NWPU VHR-10	Remote Sensing	Urban infrastructure classification; civilian applications (building damage assessment, transport monitoring)
Xu et al. (2023)	ISPRS Potsdam, Vaihingen, INRIA	Remote Sensing	HA U-Net; building extraction
Zhong et al. (2023)	In-house aerial/satellite imagery	Remote Sensing	Aerial object recognition for civilian infrastructure mapping

Study	Dataset(s) Used	Domain	Notes
Li et al. (2022)	CrackForest, CFD, in-house pavement	Infrastructure	SoUNet; pavement crack identification
Gui et al. (2023)	LS-3D-1/2/3/4, LCMS-I, ESAR, 2D-I	Infrastructure	PCDM-HED; cross-scene 3D crack detection
Xu et al. (2022)	Bridge_Crack_Image_Data, CFD	Infrastructure	CDFFHNet; channel attention; crack seg.
Guo et al. (2023)	In-house railway track imagery	Infrastructure	DeepLabV3 + HED railway defect detection
Chen et al. (2018)	In-house cement fragment imagery	Industrial	Automated cement fragment segmentation
Li & Zhang (2016)	In-house manufactured parts	Industrial	High-precision contour location
Yu et al. (2022)	In-house fabric/cloth imagery	Industrial	Cloth pattern cutting with HED
Li et al. (2021)	In-house workpiece dimension dataset	Industrial	Online visual measurement
Zhao et al. (2023)	In-house plant root images	Agricultural	Root contour extraction; edge devices
García-Peraza-Herrera et al. (2017)	MICCAI 2015 EndoVis (Da Vinci DVR)	Medical / Surgical	ToolNet; robotic surgical tool segmentation
Shoukat et al. (2023)	In-house watermarked image dataset	Security	Digital image watermarking detection
Wang et al. (2018)	BSDS500, PASCAL VOC, ADE20K	General CV	HED-CRF; semantic seg. + edge fusion
Melikyan & Melikyan (2024)	BIPED, BSDS500, in-house SEM images	General CV / Industrial	LDC; Pixel Difference Convolution
Ji et al. (2023)	In-house fingertip gesture dataset	HCI / Gesture	HED + HOG for fingertip detection

Note: BSDS500 = Berkeley Segmentation Data Set 500 Arbeláez et al. (2011); NYUD = NYU Depth Dataset Silberman et al. (2012); BIPED = Barcelona Images for Perceptual Edge Detection; MBD = Multicue Dataset for Boundary Detection; CID = Contour Image Dataset; DCD = DexiNed Complementary Dataset; PROMISE12 = Prostate MR Image Segmentation 2012; LiTS = Liver Tumour Segmentation Challenge; OAI = Osteoarthritis Initiative; CFD = Crack Forest Dataset; DOTA = Detection in Optical Remote Sensing Images (civilian applications) “in-house” = proprietary dataset, not publicly available.

In **medical imaging**, optimised HED assists with organ segmentation (prostate on PROMISE12 Cheng et al. (2017); Sumathipala et al. (2018); liver on LiTS Mu-

lay et al. (2023)), retinal vessel tracing on DRIVE/STARE/CHASEDB1 Jain & Dwivedi (2024), and osteoarthritis grading on OAI knee X-rays Yang et al. (2022). Domain constraints include heterogeneous intensity distributions across MRI acquisition protocols, class imbalances between foreground tissue and background, and the near absence of large publicly labelled datasets.

For **civil infrastructure**, HED detects fine linear cracks in both 2D (Crack-Forest, CFD) and 3D laser-scanning datasets Li et al. (2022); Gui et al. (2023). Concrete surface evaluation used Bridge_Crack_Image_Data and CFD Xu et al. (2022), and railway defect detection used in-house track imagery Guo et al. (2023).

In **remote sensing**, projects monitor Antarctic coastlines Heidler et al. (2021), identify building footprints on ISPRS Potsdam, SpaceNet, and WHU Building datasets Jung et al. (2021); Xu et al. (2023), support civilian infrastructure classification on DOTA and NWPU VHR-10 Tiwari et al. (2024), and detect sea-sky horizon lines in maritime imagery Chen et al. (2024).

In **industrial and agricultural** settings, HED is deployed for workpiece measurement, fabric pattern cutting, and plant root extraction, all using in-house proprietary datasets Li et al. (2021); Yu et al. (2022); Zhao et al. (2023). The near-universal reliance on in-house data in these sectors underscores the domain adaptation challenge identified in RQ5.

7.3. RQ3: Preprocessing and Postprocessing Pipeline Components

HED optimisation depends heavily on auxiliary pipeline stages to handle data noise and refine the outputs.

Preprocessing techniques include median filtering and IQR scaling to normalise intensities in MRI and X-ray scans Cheng et al. (2017); Yang et al. (2022). For the 3D pavement datasets used in the PCDM-HED, a Profile Component Decomposition Model removes vehicle fluctuations and pavement deformation before passing depth maps to the network Gui et al. (2023). Otsu-Canny operators are frequently used to provide initial rough edge maps that HED then refines Wang et al. (2022).

Postprocessing techniques, such as non-maximum suppression (NMS) and double-thresholding, are standard for thinning HED probability maps into single-pixel edges Li et al. (2021). Advanced pipelines use meta-heuristic algorithms (SSO, BWO) to refine the edges in colour images Heidler et al. (2021). Other steps include skeletonisation for maritime horizons and mRMR feature selection to reduce dimensionality before the final classification on DOTA and NWPU VHR-10 Tiwari et al. (2024).

7.4. RQ4: Performance Metrics and Comparative Analysis

The primary metrics across studies are the Optimal Dataset Scale (ODS), Optimal Image Scale (OIS), Average Precision (AP), and F1-score Wang et al. (2022). Medical and industrial tasks also use AUC, PSNR, and SSIM Sumathipala et al. (2018).

7.4.1. Descriptive Comparative Analysis of Performance Improvements

ODS scores were extracted from 18 of the 42 studies that reported BSDS500 benchmark results, enabling a descriptive comparative analysis. A formal statistical meta-analysis was not possible. The synthesis below is a structured narrative comparison, not a causal inference. Figure 3 presents the distribution of ODS improvements over baseline HED (0.788), categorised by optimisation type. Only 7 of these 18 studies (38.9%) reported confidence intervals or statistical significance testing, which limits the strength of cross-study comparisons.

The key observations from this descriptive comparison are as follows.

1. **Meta-heuristic approaches show highest reported improvement, with important caveats:** Three studies employing SSO or BWO refinement on BSDS500 yield a median ODS gain of 2.6 pp, with the best-performing study Zhang et al. (2024) reaching 0.806 ODS (+1.8 pp over baseline). However, these three studies are the only evidence base for this claim. None of the studies reported confidence intervals or statistical significance testing, and the computational cost of meta-heuristic refinement (convergence iterations, sensitivity to initialisation) was not documented




Figure 3: Distribution of ODS improvements over baseline HED (0.788) by optimisation category. Box plots show median, IQR, and outliers; numbers indicate study count per category.

in any of the three studies. Therefore, the meta-heuristic improvement claim should be treated as a promising finding from a small sample rather than a definitive performance advantage. A formal cost-benefit analysis quantifying GPU hours against ODS gain was not possible using the available data. Replication on larger study sets is a clear priority for future studies.

2. **Attention mechanisms show modest but consistent gains:** Channel and spatial attention yield median improvements of +1.4 pp ODS with lower variance than other categories, suggesting reliable if bounded benefits.
3. **Backbone replacement involves accuracy-computation trade-offs:** ResNet-50 improvements (median +1.8 pp) come with an 18% FLOP increase (Table 7), which raises questions about cost-effectiveness for resource-constrained deployment.

4. **Reporting heterogeneity prevents strong inference:** Only 7 of 18 studies reported confidence intervals or significance tests. The hardware platforms varied widely. Standardised benchmarking protocols are urgently needed in this field.

Medical imaging studies consistently report higher relative improvements (AUC gains of 3–7%) because baseline HED struggles with heterogeneous intensity distributions in PROMISE12 and multi-parametric MRI Cheng et al. (2017); Sumathipala et al. (2018). Remote sensing applications show variable results (ODS range: 0.712–0.834), reflecting dataset diversity across ISPRS Potsdam, SpaceNet, DOTA, and in-house datasets Heidler et al. (2021); Chen et al. (2024).

7.5. RQ5: Technical Challenges and Reproducibility Gaps

Despite recent advances, several challenges persist. Standard HED loses target contour information in low-resolution high-level feature maps, as demonstrated in BSDS500 evaluations Wang et al. (2022). Real-time deployment on edge devices is difficult because complex hybrid models increase processing time and memory consumption, particularly for plant root extraction and aerial imagery applications Zhong et al. (2023).

Many domains lack large-scale, manually labelled, public ground-truth datasets, making deep learning training difficult and limiting reproducibility. As shown in Table 9, 26 of the 42 studies (61.9%) used in-house proprietary datasets, effectively preventing independent replication.

Models also struggle with over-detection or missed features in scenes with complex backgrounds or extreme lighting, as reported in BSDS500 ablations and industrial in-house datasets Wang et al. (2022).

7.6. RQ6: Hybrid Approaches and Multi-Objective Optimisation

Hybrid approaches have transformed HED from a pure edge detector to a multitask feature extractor. Integrating HED with U-Net on Antarctic coastline data or with Mask R-CNN on LiTS liver data allows simultaneous semantic

segmentation and edge detection, helping models learn class boundaries while maintaining global semantic context Heidler et al. (2021).

Multi-objective loss functions (BCE + Boundary Loss + Texture Suppression) balance precise edge localisation against background noise suppression, as demonstrated on BIPED and BSDS500 Heidler et al. (2021); Melikyan & Melikyan (2024). Meta-heuristic approaches (SSO, BWO) refine HED outputs by minimising structural differences across multiple colour layers on the BSDS500, making edges more consistent across spectral channels Zhang et al. (2024).

7.7. Temporal Evolution of Optimisation Strategies (2015–2025)

Analysis of publication dates reveals distinct phases in HED optimisation research (Figure 4):




Figure 4: Temporal distribution of HED optimisation studies by theme (2015–2025), showing evolution from architectural modifications toward hybrid and meta-heuristic approaches.

- **2015–2017 (Foundation Phase):** Early work focused on architectural modifications (ResNet backbones, short connections) and established base-

line improvements on BSDS500 and NYUD. The average quality scores for this period were 72.3%, reflecting strong methodological foundations.

- **2018–2020 (Attention Era):** Integration of channel and spatial attention mechanisms dominated, driven by broader computer vision successes, primarily evaluated on BSDS500. Average quality scores declined to 65.8% with reduced code availability (15% vs. 40% in the Foundation Phase).
- **2021–2023 (Domain Specialisation):** The field shifted toward domain-specific adaptations: medical on PROMISE12 and LiTS; remote sensing on ISPRS and SpaceNet; industrial on in-house datasets. Quality scores recovered to 69.4% with increased dataset diversity but persistent reproducibility gaps.
- **2024–2025 (Hybrid Intelligence):** Meta-heuristic-neuro hybrid approaches (SSO, BWO, mRMR) emerged, evaluated on BIPED, BSDS500, DRIVE/STARE/CHASEDB1, and DOTA. These studies showed the highest quality scores (74.1%) and improved code availability (28%), suggesting maturing methodological standards.

7.7.1. *Emerging Paradigms: 2024–2025 Breakthroughs Beyond HED Optimisation*

While this review focuses specifically on HED optimisation, the broader edge detection field has seen architectural shifts in 2024–2025 that practitioners should be aware of when selecting a baseline.

Diffusion-based edge detection (DiffusionEdge, AAAI 2024) applies diffusion probabilistic models in latent space, reporting substantial relative gains on NYUDv2 compared to CNN-based methods. This approach fundamentally differs from HED optimisation in that it replaces the deterministic encoder-decoder paradigm with a generative inference process at the cost of a substantially higher inference time.

Transformer-based approaches are maturing: ECT (2023) introduces learned cause tokens for semantic boundary detection; EdgeNAT (2025) pro-

poses neighbourhood attention for computationally efficient transformer-based edge detection without the quadratic attention scaling of standard Vision Transformers.

Self-supervised and domain-generalisation approaches (SuperEdge, 2024) address the domain shift problem that HED-based models consistently struggle with, achieving better generalisation to novel domains without extensive labelled data.

Frequency-domain fusion (FSFD, 2024, *Pattern Recognition*) combines spatial and frequency features to handle low-contrast scenes, which remains a known weakness of HED across BSDS500 evaluations.

These developments do not invalidate the HED optimisation findings of this review. HED-based architectures remain practical for deployment contexts in which model interpretability, architectural transparency, and lightweight inference are priorities. However, for tasks where training data are abundant and accuracy on general benchmarks is the primary goal, transformer- and diffusion-based approaches are now competitive alternatives that any future HED optimisation study should benchmark against.

8. Conclusion

This systematic literature review synthesises evidence from 42 primary studies published between 2015 and 2025 on HED optimisation techniques. Using the PICOC framework and PRISMA 2020 guidelines, we established a reproducible knowledge base addressing critical gaps in the fragmented HED optimisation literature.

8.1. Summary of Key Findings

First, HED optimisation has progressed through distinct phases: backbone replacement and architectural modification (2015–2018), attention mechanism integration (2019–2021), and current hybrid approaches combining deep learning with meta-heuristic search and multi-task learning (2022–2025). The field

has moved from generic accuracy improvements toward deployment-aware optimisation that addresses specific computational and domain constraints.

Second, no single optimisation strategy dominates across all evaluation criteria. Backbone replacement (ResNet-50, ConvNeXt) delivers consistent accuracy gains (+1.8 pp ODS median on BSDS500) but increases the computational costs. Lightweight architectures (LDC, DexiNed) achieve 3–4× speed improvements with modest accuracy trade-offs. Attention mechanisms offer reliable gains (+1.4 pp ODS) with minimal overhead costs. Meta-heuristic hybrids showed the highest reported improvement potential (+2.6 pp ODS) across three contributing studies, although this finding is preliminary and requires replication at scale.

Third, reproducibility deficits are severe. Only 2.4% of the assessed studies achieved excellent quality ratings ($\geq 90\%$), code availability was low (19% of studies), and 61.9% of domain-specific studies used proprietary in-house datasets (Table 9). This opacity undermines independent verification, fair comparison, and cumulative knowledge building.

Fourth, domain-specific optimisation is a critical success factor. Generic HED modifications generalise poorly across medical imaging, remote sensing, and industrial inspection because of distinct intensity distributions, noise characteristics, and boundary definitions. Effective deployment requires domain-aware preprocessing and postprocessing pipelines that are currently undocumented in the literature.

Fifth, environmental and deployment costs remain invisible. Despite stated goals of "efficiency" and "lightweight" deployment, only 23.8% of studies reported inference time on standardised hardware, and none quantified the carbon footprint or energy consumption.

8.2. Recommendations for Practitioners

- **Resource-constrained deployment:** DexiNed or LDC-HED variants over attention-enhanced or hybrid models. Consider INT8 quantisation for an additional 2–4× speedup with <1% accuracy loss.

- **High-accuracy requirements:** ResNet-50 backbones with multi-scale fusion for the best accuracy-computation balance. Complex meta-heuristic postprocessing in latency-sensitive applications should be avoided.
- **Domain-specific applications:** Budget for domain-specific preprocessing pipeline development. Generic HED optimisations require 15–30% accuracy adaptation for medical (PROMISE12, LiTS) and remote sensing (ISPRS, SpaceNet) tasks.
- **Dataset standardisation:** Prefer studies benchmarked on publicly available datasets (BSDS500, BIPED, DRIVE, CrackForest, CFD). Make the code and dataset availability a precondition for integration into production pipelines.
- **Maintenance and monitoring:** Implement automated regression testing for edge detection models, as HED optimisations show sensitivity to input distribution shifts not captured by standard accuracy metrics.

8.3. Limitations of this Literature Review

This review has several constraints that bound the extent to which the conclusions can be generalised.

Publication bias likely favours positive results, potentially inflating the reported optimisation benefits. Our quality assessment, while systematic, involved subjective judgement in criterion scoring, which was mitigated through author consensus and explicit weighting rationale. The rapidly evolving field means that post-2025 developments (transformer-based edge detection, diffusion models for boundary refinement) are not fully captured, although Section 7.7.1 provides a brief orientation. The HED-specific focus excludes direct comparative analysis with CASENet, RCF, and BDCN, which may offer superior optimisation potential for specific use cases. The predominance of in-house datasets across 26 of 42 studies limits the generalisability of meta-analytic comparisons and constrains cross-study conclusions. Finally, the three-study evidence base

for meta-heuristic approaches (Theme 5) is too small to draw strong conclusions about their performance advantages.

8.4. *Broader Impact*

Only 10 of 42 studies (23.8%) reported inference time on standardised hardware, and none reported energy consumption or carbon footprints. This is a significant gap, given that many HED optimisation studies frame their contributions in terms of computational efficiency. Future work should report CO₂ equivalent per inference alongside accuracy metrics, particularly for meta-heuristic and neural architecture search methods, which can require substantial training computations. The lightweight architectures identified in this review (DexiNed and LDC-HED) are the most environmentally responsible choices for large-scale deployment. However, this must be quantified in absolute terms rather than relative speedup ratios.

Edge detection is an enabling technology with dual-use potential. In civilian contexts covered by this review, such as medical imaging, infrastructure inspection, and remote sensing for urban planning, the applications are broadly advantageous. However, boundary detection is also a component of surveillance, facial recognition, and autonomous weapon systems. Researchers publishing HED optimisation work are encouraged to include explicit intended-use statements and consider whether their chosen datasets and evaluation benchmarks carry implicit application assumptions. The use of remote sensing datasets such as DOTA for urban infrastructure classification in this review is framed explicitly within civilian applications (building damage assessment, transportation monitoring, and disaster response).

The 61.9% proprietary dataset rate identified in this review creates a barrier not only to reproducibility but also to participation by researchers at institutions without access to specialised data acquisition equipment. Centralised, publicly labelled benchmark datasets for high-value domains (crack detection, medical boundary segmentation) would substantially lower the entry barrier. The System Designer’s Selection Guide in Section 6 is intended to be usable by

practitioners regardless of institutional resources.

References

- Acuna, D., & Fidler, S. (2022). Towards a general understanding of edge detection. [arXiv:2205.01895](#) arXiv preprint; arXiv ID may require verification by authors.
- Arbeláez, P., Maire, M., Fowlkes, C., & Malik, J. (2011). Contour detection and hierarchical image segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, *33*, 898–916. doi:doi:10.1109/TPAMI.2010.161.
- Beaini, D., Achiche, S., Duperré, A., & Raison, M. (2020). Deep green function convolution for improving saliency. *Neural Computing and Applications*, *34*, 10587–10596. doi:doi:10.1007/s00521-020-05205-6.
- Chen, H., Jin, Y., Li, G., & Chu, B. (2018). Automated cement fragment image segmentation and distribution estimation via a holistically-nested convolutional network and morphological analysis. *Powder Technology*, *339*, 306–313. doi:doi:10.1016/j.powtec.2018.08.014.
- Chen, Q., Wang, Z., Zhang, D., & Zhang, X. (2024). Innovation in sea-sky line detection: Fusion of ConvNeXt classification, HED, and canny edge detection. In *Proceedings of the 2024 International Conference on Intelligent Perception and Pattern Recognition IPPR '24* (pp. 1–8). New York, NY, USA: Association for Computing Machinery. doi:doi:10.1145/3674393.3674400.
- Cheng, R., Roth, H. R., Lay, N., Lu, L., Turkbey, B., Gandler, W., McCreedy, E. S., Pohida, T., Pinto, P. A., Choyke, P., McAuliffe, M. J., & Summers, R. M. (2017). Automatic magnetic resonance prostate segmentation by deep learning with holistically nested networks. In *Proceedings of SPIE Medical Imaging 2017: Image-Guided Procedures, Robotic Interventions, and Modeling* (p. 101351I). SPIE volume 10135. doi:doi:10.1117/12.2254621.
- Dalkey, N., & Helmer, O. (1963). An experimental application of the Delphi method to the use of experts. *Management Science*, *9*, 458–467.

- Gamal, S., Atef, H., Youssef, D., Ismail, T., & El-Azab, J. (2023). Early breast cancer screening from thermography via deep pre-trained edge detection with extreme gradient boosting. In *2023 Intelligent Methods, Systems, and Applications (IMSA)* (pp. 430–433). IEEE. doi:doi:10.1109/IMSA58816.2023.10340872.
- García-Peraza-Herrera, L. C., Li, W., Fidon, L., Gruijthuijsen, C., Devreker, A., Attilakos, G., Deprest, J., Vander Poorten, E., Stoyanov, D., Vercauteren, T., & Ourselin, S. (2017). ToolNet: Holistically-nested real-time segmentation of robotic surgical tools. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)* (pp. 5717–5722). IEEE. doi:doi:10.1109/IROS.2017.8206462.
- Gui, R., Sun, Q., Wu, W., Zhang, D., & Li, Q. (2023). Transfer learning for cross-scene 3d pavement crack detection based on enhanced deep edge features. *Engineering Applications of Artificial Intelligence*, *123*, 106452. doi:doi:10.1016/j.engappai.2023.106452.
- Guo, R., Liu, R., Li, N., & Liu, W. (2023). DV3+HED+: A DCNNs-based framework to monitor temporary works and ESAs in railway construction projects using VHR satellite images. *Remote Sensing*, *15*, 1837. doi:doi:10.3390/rs15071837.
- Heidler, K., Mou, L., Baumhoer, C., Dietz, A., & Zhu, X. X. (2021). HED-UNet: a multi-scale framework for simultaneous segmentation and edge detection. In *2021 IEEE International Geoscience and Remote Sensing Symposium (IGARSS)* (pp. 3037–3040). IEEE. doi:doi:10.1109/IGARSS47720.2021.9553493.
- Hou, Q., Cheng, M.-M., Hu, X., Borji, A., Tu, Z., & Torr, P. H. S. (2016). Deeply supervised salient object detection with short connections. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 3203–3212). IEEE. doi:doi:10.1109/CVPR.2016.349.

- Hu, A. (2025). A mathematical survey of deep edge detection: From convolution to attention. *[Journal Name]*, *[volume]*, [pages].
- Hu, G., & Saeli, C. (2021). Scale-invariant salient edge detection. In *2021 IEEE International Conference on Image Processing (ICIP)* (pp. 564–568). IEEE. doi:doi:10.1109/ICIP42928.2021.9506491.
- Jain, A., & Dwivedi, R. K. (2024). Hybrid segmentation-based agricultural leaf disease detection (Hy-SALDD) using black widow optimisation for feature selection and Bayesian-optimised SVM classification. *Artificial Intelligence Review*, *57*, 168. doi:doi:10.1007/s10462-024-10845-5.
- Ji, D., Qian, J., Yu, J., Kurihara, T., & Zhan, S. (2020). Automatic prostate segmentation on MR images using enhanced holistically-nested networks. *Journal of Healthcare Engineering*, *2020*, 8842066. doi:doi:10.1155/2020/8842066.
- Ji, J., Lu, X., Luo, M., Yin, M., Miao, Q., & Liu, X. (2021). Parallel fully convolutional network for semantic segmentation. *IEEE Access*, *9*, 673–682. doi:doi:10.1109/ACCESS.2020.3042254.
- Ji, P., Feng, J., Ma, F., Wang, X., & Li, C. (2023). Fingertip detection algorithm based on maximum discrimination HOG feature in complex background. *IEEE Access*, *11*, 3160–3173. doi:doi:10.1109/ACCESS.2023.3234526.
- Jiao, Y., Wang, X., Wang, W., & Li, S. (2022). Image semantic segmentation fusion of edge detection and AFF attention mechanism. *Applied Sciences*, *12*, 11248. doi:doi:10.3390/app122111248.
- Jung, H., Choi, H.-S., & Kang, M. (2021). Boundary enhancement semantic segmentation for building extraction from remote sensed image. *IEEE Transactions on Geoscience and Remote Sensing*, *60*, 1–12. doi:doi:10.1109/TGRS.2021.3132567.

- Kitchenham, B., & Charters, S. (2007). *Guidelines for Performing Systematic Literature Reviews in Software Engineering*. Technical Report EBSE-2007-01 EBSE (Evidence-Based Software Engineering).
- Koo, T. K., & Li, M. Y. (2016). A guideline of selecting and reporting intra-class correlation coefficients for reliability research. *Journal of Chiropractic Medicine*, 15, 155–163.
- Landis, J. R., & Koch, G. G. (1977). The measurement of observer agreement for categorical data. *Biometrics*, 33, 159–174.
- Li, J., & Zhang, Q. (2016). A method for contour location of industrial objects based on CNN. In *2016 31st Youth Academic Annual Conference of Chinese Association of Automation (YAC)* (pp. 247–250). IEEE. doi:doi:10.1109/YAC.2016.7804927.
- Li, P., Xia, H., Zhou, B., Yan, F., & Guo, R. (2022). A method to improve the accuracy of pavement crack identification by combining a semantic segmentation and edge detection model. *Applied Sciences*, 12, 4714. doi:doi:10.3390/app12094714.
- Li, W., & He, M. (2022). Improved edge detection model based on HED. *Journal of Physics: Conference Series*, 2386, 012019. doi:doi:10.1088/1742-6596/2386/1/012019.
- Li, X., Yang, Y., Ye, Y., Ma, S., & Hu, T. (2021). An online visual measurement method for workpiece dimension based on deep learning. *Measurement*, 185, 110032. doi:doi:10.1016/j.measurement.2021.110032.
- Lou, L., & Zang, S. (2023). Research on edge detection method based on improved HED network. *Journal of Physics: Conference Series*, 2469, 012020. doi:doi:10.1088/1742-6596/2469/1/012020.
- Melikyan, S., & Melikian, G. (2024). Advanced edge detection for noisy SEM images: A hybrid approach for enhanced precision. In *2024*

- IEEE East-West Design & Test Symposium (EWDTS)* (pp. 1–4). IEEE. doi:doi:10.1109/EWDTS63723.2024.10873599.
- Mubashar, A. (2022). Analysis of dataset limitations and deep neural network techniques in edge detection. *[Journal Name]*, *[volume]*, [pages].
- Mulay, S., G, D., S, J., Ram, K., & Sivaprakasam, M. (2023). Liver segmentation from multimodal images using HED-Mask R-CNN. In *IEEE International Symposium on Biomedical Imaging (ISBI)* (pp. 1–5). IEEE. doi:doi:10.1109/ISBI53787.2023.10230501.
- Nogues, I., Lu, L., Wang, X., Roth, H., Bertasius, G., Lay, N., Shi, J., Tsehay, Y., & Summers, R. M. (2016). Automatic lymph node cluster segmentation using holistically-nested neural networks and structured optimization in CT images. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)* (pp. 388–397). Springer. doi:doi:10.1007/978-3-319-46723-8_45.
- Page, M. J., McKenzie, J. E., Bossuyt, P. M., Boutron, I., Hoffmann, T. C., Mulrow, C. D., Shamseer, L., Tetzlaff, J. M., Akl, E. A., Brennan, S. E. et al. (2021). The PRISMA 2020 statement: An updated guideline for reporting systematic reviews. *BMJ*, *372*, n71. doi:doi:10.1136/bmj.n71.
- Pineau, J., Vincent-Lamarre, P., Sinha, K., Larivière, V., Beygelzimer, A., d’Alché Buc, F., Fox, E. B., & Larochelle, H. (2021). Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program. *Journal of Machine Learning Research*, *22*, 1–20.
- Shoukat, M. Z., Su, Z., & Ali, J. (2023). Security analysis of digital image watermarking using deep learning inspired LSB and chaotic S-Box in cyber security. *Multimedia Tools and Applications*, *82*, 23687–23713. doi:doi:10.1007/s11042-023-14656-3.
- Silberman, N., Hoiem, D., Kohli, P., & Fergus, R. (2012). Indoor segmentation and support inference from RGBD images. In *European Conference on*

- Computer Vision (ECCV)* (pp. 746–760). Springer. doi:doi:10.1007/978-3-642-33715-4_54.
- Soria, X., Riba, E., & Sappa, A. (2020). Dense extreme inception network: Towards a robust CNN model for edge detection. In *IEEE Winter Conference on Applications of Computer Vision (WACV)* (pp. 1912–1921). Snowmass Village, CO, USA: IEEE. Computer Vision Center - Universitat Autònoma de Barcelona, Barcelona, Spain; Escuela Superior Politécnica del Litoral, Guayaquil, Ecuador.
- Soria, X., & Sappa, A. (2021). Improving edge detection in RGB images by adding NIR channel. *Sensors*, *21*, 5928. doi:doi:10.3390/s21175928.
- Sreelakshmi, V., & Vardhan, D. V. (2024). Multi-magnification mask R-CNN and HED-based segmentation with hybrid DTCWT-MSER feature extraction and optimised dual-classifiers for automated colon biopsy malignancy grading. *Biomedical Signal Processing and Control*, *86*, 105171. doi:doi:10.1016/j.bspc.2023.105171.
- Sumathipala, Y., Lay, N., Turkbey, B., Smith, C., Choyke, P. L., & Summers, R. M. (2018). Prostate cancer detection from multi-institution multiparametric MRIs using deep convolutional neural networks. *Journal of Medical Imaging*, *5*, 044503. doi:doi:10.1117/1.JMI.5.4.044503.
- Sun, A. (2022). A survey of traditional and deep learning edge detectors. [*Journal Name*], [*volume*], [pages].
- Tiwari, S., Somesh, V., Joshi, R. V., & Sharma, G. P. (2024). A multi-scale hybrid segmentation and mRMR feature selection approach for remote sensing object detection. *Multimedia Tools and Applications*, *83*, 1–22. doi:doi:10.1007/s11042-023-15927-1.
- Tseng, S.-H., & Sun, W.-H. (2023). Sea-land segmentation using HED-UNet for coastline monitoring. *Remote Sensing*, *15*, 3946. doi:doi:10.3390/rs15163946.

- Wang, B., Guo, Y., & Zhu, Q. (2018). Research on image semantic segmentation algorithm based on fully convolutional HED-CRF. In *2018 Chinese Automation Congress (CAC)* (pp. 3055–3058). IEEE. doi:doi:10.1109/CAC.2018.8623528.
- Wang, Z., Li, K., Wang, X., & Lee, A. (2022). An image edge detection algorithm based on multi-feature fusion. *Computers, Materials & Continua*, *73*, 4995–5009. doi:doi:10.32604/cmc.2022.029650.
- Wen, Q., Xu, F., Lu, M., & Yong, J.-H. (2017). Real-time 3d eyelids tracking from semantic edges. *ACM Transactions on Graphics (TOG)*, *36*, 1–11. doi:doi:10.1145/3130800.3130891.
- Xie, S., & Tu, Z. (2015). Holistically-nested edge detection. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)* (pp. 1395–1403). IEEE.
- Xu, L., Liu, Y., Yang, P., Chen, H., Zhang, H., Wang, D., & Zhang, X. (2023). HA_U-Net: Improved model for building extraction from high-resolution remote sensing imagery. *IEEE Geoscience and Remote Sensing Letters*, *20*, 8003105. doi:doi:10.1109/LGRS.2023.3283105.
- Xu, S., Hao, M., Liu, G., Meng, Y., Han, J., & Shi, Y. (2022). Concrete crack segmentation based on convolution–deconvolution feature fusion with holistically nested networks. *Structural Control and Health Monitoring*, *29*, e2965. doi:doi:10.1002/stc.2965.
- Yang, S., Yu, Z., He, J., & Huang, K. (2022). A method for the grading and recognition of osteoarthritis X-ray images based on the MED-HED detection. *BioMedical Engineering OnLine*, *21*, 56. doi:doi:10.1186/s12938-022-01012-2.
- Yu, N., Zhang, Z., Xu, Q., Firdaous, E., & Lin, J. (2022). An improved method for cloth pattern cutting based on holistically-nested edge detection. *Textile Research Journal*, *92*, 3335–3348. doi:doi:10.1177/00405175221092751.

- Yuan, Y., Chen, W., Tang, J., & Yang, J. (2023). A high-precision edge detection algorithm based on improved sobel operator-assisted HED. *Sensors*, *23*, 6883. doi:doi:10.3390/s23156883.
- Zhang, J., Wang, W., & Wang, J. (2024). Edge detection in colored images using parallel CNNs and social spider optimization. *Electronics*, *13*, 3540. doi:doi:10.3390/electronics13173540.
- Zhao, P., Tan, T. F., & Wang, L. (2023). Machine-learning-algorithm-based plant root contour extraction in edge devices. *Computers and Electronics in Agriculture*, *211*, 108023. doi:doi:10.1016/j.compag.2023.108023.
- Zhong, W., Jiang, Y., & Zhang, X. (2023). Research on automatic recognition method of artificial ground target based on improved HED. *Applied Sciences*, *13*, 3163. doi:doi:10.3390/app13053163.
- Zhou, W., Li, X., & Reynolds, M. (2021). Edge detection: A systematic review and comparative evaluation. *Neurocomputing*, *433*, 1–26. doi:doi:10.1016/j.neucom.2021.03.136.